



SESSÃO 1 - CONFIGURAÇÃO DO AMBIENTE

AULA 1.1 WSL

Links:

Como instalar o Linux no Windows com o WSL:

<https://learn.microsoft.com/pt-br/windows/wsl/install>

Comandos:

Powershell:

```
wsl --version  
wsl --set-default-version 2  
wsl --list --online  
wsl --install -d Ubuntu-24.04  
python3 --version
```

AULA 1.2 Conda

Links:

Installing Miniconda: <https://www.anaconda.com/docs/getting-started/miniconda/install>

Comandos:

WSL:

```
cd ~  
wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh  
bash Miniconda3-latest-Linux-x86_64.sh  
source ~/.bashrc  
conda config --set auto_activate_base false  
rm Miniconda3-latest-Linux-x86_64.sh  
conda --version
```

AULA 1.3 Git

Links:

Downloads: <https://git-scm.com/downloads>

Comandos:

WSL:

```
sudo apt-get install git  
git --version
```

AULA 1.4 Docker

Links:

Downloads: <https://docs.docker.com/engine/install>

Comandos:

WSL:

```
sudo apt-get update  
sudo apt-get install ca-certificates curl  
sudo install -m 0755 -d /etc/apt/keyrings  
sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o  
/etc/apt/keyrings/docker.asc  
sudo chmod a+r /etc/apt/keyrings/docker.asc  
echo \  
  "deb [arch=$(dpkg --print-architecture)  
signed-by=/etc/apt/keyrings/docker.asc]  
https://download.docker.com/linux/ubuntu \  
  $(. /etc/os-release && echo "${UBUNTU_CODENAME:-$VERSION_CODENAME}")  
stable" | \  
  sudo tee /etc/apt/sources.list.d/docker.list > /dev/null  
sudo apt-get update  
sudo apt-get install docker-ce docker-ce-cli containerd.io docker-  
buildx-plugin docker-compose-plugin  
  
sudo usermod -aG docker $USER
```

```
newgrp docker
```

```
docker --version
```

AULA 1.5 VSCode

Links:

VSCode: <https://code.visualstudio.com>

AULA 1.6 GitHub, DagsHub, DockerHub

Links:

GitHub: <https://github.com/?locale=pt-br>

DagsHub: <https://dagshub.com>

DockerHub: <https://hub.docker.com>



SESSÃO 2 - PROJETO BASE

AULA 2.1 Estrutura do projeto

Links:

Repositório no GitHub: https://github.com/jonesgranatyr/mlops_project

AULA 2.2 Preparo local

Comandos:

Terminal (WSL):

```
cd ~  
mkdir PythonProjects  
ls -la  
rm -rf .git
```

```
pip install -e .
git init
git config --global user.name <your-name>
git config --global user.email <your-email>
```



Código:

.gitignore:

```
ml_classifier.egg-info
```



AULA 2.3 Pacote src

Sem material auxiliar.



AULA 2.4 Módulo load_data



Debug console:

```
data.isna().sum()/len(data)
```

Terminal:

```
python -m src.data_loading.load_data
```



Código:

.gitignore:

```
data/raw/raw.csv
```

AULA 2.5 Módulo preprocess_data

Comandos:

Debug console:

```
train_features.isna().sum()  
train_features_processed.isna().sum()
```

Terminal:

```
python -m src.data_preprocessing.preprocess_data
```

Código:

.gitignore:

```
data/preprocessed/train_preprocessed.csv  
data/preprocessed/test_preprocessed.csv  
artifacts\[features\]_mean_imputer.joblib
```

AULA 2.6 Módulo engineer_features

Comandos:

Debug console:

```
train_preprocessed  
train_processed
```

Terminal:

```
python -m src.feature_engineering.engineer_features
```

Código:

.gitignore:

```
data/processed/train_processed.csv  
data/processed/test_rocessed.csv  
artifacts\[features\]_scaler.joblib
```

AULA 2.7 Módulo train_model

Comandos:

Debug console:

params

Terminal:

```
python -m src.model_training.train_model
```

Código:

.gitignore:

```
artifacts/\\[target\\]_one_hot_encoder.joblib
models/model.keras
metrics/training.json
```

AULA 2.8 Módulo evaluate_model

Comandos:

Debug console:

y_pred_proba
y_pred
report
cm

Terminal:

```
python -m src.model_evaluation.evaluate_model
```



Código:

.gitignore:

```
metrics/evaluation.json
```



AULA 2.9 Pacote app



.gitignore:

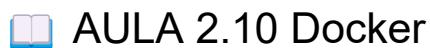
```
# Python
ml_classifier.egg-info
__pycache__

# Data
data/**/*.csv

# Artifacts
artifacts/*.joblib

# Models
models/*.keras

# Metrics
metrics/*.json
```



AULA 2.10 Docker



Terminal:

```
docker build -t ml-classifier .
docker image list
docker run -p 5001:5001 ml-classifier
```

```
docker ps  
docker container stop <container_id>
```

AULA 2.11 Repositório remoto - GitHub

Comandos:

Terminal:

```
git status  
git add .  
git commit -m "first commit"  
git remote add origin <remote-url>  
git branch -M main  
git push -u origin main  
git push
```

SESSÃO 3 - DVC

AULA 3.1 Arquivos de configuração

Comandos:

Terminal:

```
pip install -e .  
dvc init
```

Código:

pyproject.toml:

```
version = "0.2"
```

```
<...>
```

```
"dvc==3.59.1",
```

```
.dockerignore:
```

```
# DVC  
.dvc
```

```
dvc.yaml:
```

```
stages:  
  load_data:  
    cmd: python -m src.data_loading.load_data  
    deps:  
      - src/data_loading/load_data.py  
    outs:  
      - data/raw/raw.csv  
  
  preprocess_data:  
    cmd: python -m src.data_preprocessing.preprocess_data  
    deps:  
      - src/data_preprocessing/preprocess_data.py  
      - data/raw/raw.csv  
    outs:  
      - data/preprocessed/train_preprocessed.csv  
      - data/preprocessed/test_preprocessed.csv  
      - artifacts/[features]_mean_imputer.joblib  
    params:  
      - preprocess_data.test_size  
      - preprocess_data.random_seed  
  
  engineer_features:  
    cmd: python -m src.feature_engineering.engineer_features  
    deps:  
      - src/feature_engineering/engineer_features.py  
      - data/preprocessed/train_preprocessed.csv  
      - data/preprocessed/test_preprocessed.csv  
    outs:  
      - data/processed/train_processed.csv
```

```
- data/processed/test_processed.csv
- artifacts/[features]_scaler.joblib

train:
cmd: python -m src.model_training.train_model
deps:
- src/model_training/train_model.py
- data/processed/train_processed.csv
outs:
- models/model.keras
- artifacts/[target]_one_hot_encoder.joblib
metrics:
- metrics/training.json
params:
- train.learning_rate
- train.hidden_layer_1_neurons
- train.hidden_layer_2_neurons
- train.dropout_rate
- train.epochs
- train.batch_size

evaluate:
cmd: python -m src.model_evaluation.evaluate_model
deps:
- src/model_evaluation/evaluate_model.py
- models/model.keras
- artifacts/[target]_one_hot_encoder.joblib
- data/processed/test_processed.csv
metrics:
- metrics/evaluation.json
```



AULA 3.2 Execução automática da pipeline



Terminal:

```
dvc repro
```

```
git add dvc.lock  
dvc config core.autostage true  
dvc pull
```

AULA 3.3 Versionamento de dados

Comandos:

Terminal:

```
dvc repro  
dvc repro load_data  
git checkout <hash_id>  
dvc checkout  
git checkout main
```

Código:

dvc.yaml:

```
always_changed: True
```

```
load_data:
```

```
import time  
  
<...>  
  
data = data.sample(frac=1, random_state=int(time.time())).reset_index(drop=True)
```

AULA 3.4 Experimentos

Comandos:

Terminal:

```
dvc exp run -S train.batch_size=16  
dvc exp run -S preprocess_data.test_size=0.25 --queue
```

```
dvc exp run -S train.dropout_rate=0.25 --queue  
dvc exp run -S train.random_seed=24 --queue  
dvc exp run --run-all  
dvc exp show  
dvc exp apply <experiment_id>
```

AULA 3.5 Repositório remoto - DagsHub

Comandos:

Terminal:

```
dvc remote add origin <dagshub-url>  
dvc remote modify origin --local auth basic  
dvc remote modify origin --local user <username>  
dvc remote modify origin --local password <token>  
dvc remote list  
dvc remote default origin  
dvc push  
dvc pull
```

SESSÃO 4 - MLFLOW

AULA 4.1 Experimento e run

Comandos:

Terminal:

```
pip install -e .  
dvc repro
```



Código:

pyproject.toml:

```
version = "0.3"  
  
<...>  
  
"mlflow==2.22.1",
```

train_model.py:

```
import mlflow  
  
<...>  
  
# Set up MLflow experiment  
mlflow.set_experiment("ml_classification")  
  
# Set up Keras autolog  
mlflow.keras.autolog()  
  
with mlflow.start_run():  
    # Log parameters to MLflow  
    mlflow.log_params(params)  
  
<...>  
  
    # Log preprocessing artifacts  
    mlflow.log_artifact("artifacts/[features]_mean_imputer.joblib")  
    mlflow.log_artifact("artifacts/[features]_scaler.joblib")  
  
<...>  
  
    # Log the encoder  
    mlflow.log_artifact("artifacts/[target]_one_hot_encoder.joblib")  
  
<...>
```

```
# # Log metrics to MLflow
# mlflow.log_metrics(metrics)
```

evaluate_model.py:

```
import os

<...>

import mlflow

<...>

# Set up MLflow experiment
mlflow.set_experiment("ml_classification")

# Get run_id for latest MLflow run
runs = mlflow.search_runs(
    experiment_ids=[os.getenv("MLFLOW_EXPERIMENT_ID")], order_by=["start_time DESC"]
)
run_id = runs.iloc[0].run_id

with mlflow.start_run(run_id=run_id):

    <...>

    # Log metrics (DVC)

    <...>

    # Log metrics (MLflow)
    mlflow.log_metrics(
        {
            "test_accuracy": report["accuracy"],
            "test_precision_weighted": report["weighted_avg"]["precision"],
            "test_recall_weighted": report["weighted_avg"]["recall"],
            "test_f1_weighted": report["weighted_avg"]["f1-score"],
        }
    )
```

)

.gitignore:

```
# MLflow  
mlruns
```

.dockerignore:

```
# MLflow  
mlruns
```

AULA 4.2 Dashboard

 Comandos:

Terminal:

```
mlflow ui  
dvc repro
```

AULA 4.3 Experimentos do DVC no MLflow

 Comandos:

Terminal:

```
pip install -e .  
mlflow ui  
dvc exp run -S preprocess_data.test_size=0.25 --queue  
dvc exp run -S train.dropout_rate=0.25 --queue  
dvc exp run -S train.random_seed=24 --queue  
dvc exp run --run-all
```

 Código:

pyproject.toml:

```
"python-dotenv==1.1.0",
```

```
.env:
```

```
MLFLOW_TRACKING_URI=http://localhost:5000
```

```
.gitignore:
```

```
# Project
```

```
.env
```

```
src/__init__.py:
```

```
from dotenv import load_dotenv

load_dotenv()
```

```
train_model.py:
```

```
# Setting MLflow if we are running a DVC experiment
is_experiment = os.getenv("DVC_EXP_NAME") is not None
extra_args = {}
if is_experiment:
    runs = mlflow.search_runs(
        experiment_ids=[os.getenv("MLFLOW_EXPERIMENT_ID")],
        filter_string="tags.dvc_exp = 'True'",
        order_by=["start_time DESC"],
    )
    if runs.empty:
        with mlflow.start_run() as parent_run:
            mlflow.set_tag("dvc_exp", True)
            parent_run_id = parent_run.info.run_id
    else:
        parent_run_id = runs.iloc[0].run_id
run_name = os.getenv("DVC_EXP_NAME")
```

```
extra_args = {  
    "parent_run_id": parent_run_id,  
    "run_name": run_name,  
    "nested": True,  
}  
  
with mlflow.start_run(**extra_args):
```

AULA 4.4 Registro de artefatos

Comandos:

Debug console:

```
latest_run  
parent_run_id
```

Código:

register_artifacts.py:

```
import logging  
import mlflow  
from mlflow.tracking import MlflowClient  
import pandas as pd  
  
logger = logging.getLogger("src.register_artifacts")  
  
client = MlflowClient()  
  
  
def get_best_run(experiment_id: str, parent_run_id: str) -> pd.Series:  
    """Get the best child run based on test accuracy for a given parent run.  
  
    Args:  
        client: MLflow client instance  
        parent_run_id: ID of the parent run
```

```
Returns:  
    The best run as a pandas Series  
"""  
  
# Get all child runs for the parent  
child_runs = client.search_runs(  
    experiment_ids=[experiment_id],  
    filter_string=f"tags.mlflow.parentRunId = '{parent_run_id}'",  
    order_by=["metrics.test_accuracy DESC"],  
    max_results=1000  
)  
# Return the run with highest test accuracy  
return child_runs[0]  
  
def register_model() -> None:  
    """Register the model that was logged during training."""  
  
    logger.info("Registering model from latest MLflow run")  
  
    # Get the experiment ID for the 'ml_classification' experiment  
    experiment_id = client.get_experiment_by_name("ml_classification").experiment_id  
  
    # Get the latest run from the experiment  
    latest_run = client.search_runs(  
        experiment_ids=[experiment_id],  
        order_by=["start_time DESC"],  
        max_results=1  
)  
    )[0]  
  
    # Check if the latest run has a parent run  
    run_id = latest_run.info.run_id  
    parent_run_id = latest_run.data.tags.get('mlflow.parentRunId')  
  
    if parent_run_id:  
        logger.info(f"Latest run has parent run ID: {parent_run_id}")  
        best_run = client.search_runs(  
            experiment_ids=[experiment_id],  
            filter_string=f"tags.mlflow.parentRunId = '{parent_run_id}'",
```

```
    order_by=[ "metrics.test_accuracy DESC" ],
    max_results=1
)[0]
run_id = best_run.info.run_id
logger.info(f"Using best run {run_id} with test_accuracy:
{best_run.data.metrics['test_accuracy']}")

# Register the model from the run
logger.info("Registering model")
try:
    client.create_registered_model("model")
except mlflow.exceptions.MlflowException:
    logger.debug("Model already exists")

model_uri = f"runs:{run_id}/model"
client.create_model_version(
    name="model",
    source=model_uri,
    run_id=run_id
)
logger.info("Registered model successfully")

def main() -> None:
    """Main function to orchestrate the model registration process."""
    register_model()
    logger.info("Model registration completed")

if __name__ == "__main__":
    main()
```

main.py:

```
import mlflow

<...>
```

```
from mlflow.tracking import MLflowClient

<...>

# Load model from registry
logger.info("Loading registered model from MLflow Model Registry")
self.model = mlflow.keras.load_model("models:/model/latest")

# Get run_id from model version metadata
client = MLflowClient()
run_id = client.get_registered_model("model").latest_versions[0].run_id

# Load related artifacts
logger.info(f"Loading artifacts from run {run_id}")
artifacts_dir = mlflow.artifacts.download_artifacts(run_id=run_id, artifact_path="")

imputer_path = os.path.join(artifacts_dir, "[features]_mean_imputer.joblib")
self.features_imputer = joblib.load(imputer_path)
scaler_path = os.path.join(artifacts_dir, "[features]_scaler.joblib")
self.features_scaler = joblib.load(scaler_path)
encoder_path = os.path.join(artifacts_dir, "[target]_one_hot_encoder.joblib")
self.target_encoder = joblib.load(encoder_path)

logger.info("Successfully loaded model and related artifacts")
```

AULA 4.5 Repositório remoto - DagsHub

Comandos:

Terminal:

```
pip install -e .
dvc repro
python -m src.register_artifacts
```



Código:

pyproject.toml:

```
"dagshub==0.5.9",
```

.env:

```
MLFLOW_TRACKING_URI=<dagshub-url>
```

src/__init__.py:

```
import dagshub

<...>

# Initialize DagsHub with credentials
dagshub.init(
    repo_owner=<your-user-name>,
    repo_name=<your-repo-name>
)
```



AULA 4.6 Docker



Links:

Criação de tokens DagsHub: <https://dagshub.com/user/settings/tokens>



Comandos:

Terminal:

```
docker build -t ml-classifier .
docker run -p 5001:5001 ml-classifier
dvc push
```



Código:

.dockerignore:

```
# Project
```

artifacts
models

app/__init__.py:

```
from dotenv import load_dotenv

load_dotenv()
```

.env:

```
DAGSHUB_USER_TOKEN=<dagshub-token>
```



SESSÃO 5 - APACHE AIRFLOW

📘 AULA 5.1 Arquivo dag

💡 Comandos:

Terminal:

```
pip install -e .
airflow db init
cd ~
cd airflow
ls -la
airflow users create --username admin --password admin --firstname
admin --lastname admin --role Admin --email admin@example.org
airflow webserver
airflow scheduler
cd ~/airflow
cat airflow.cfg
grep load_examples airflow.cfg
sed -i 's|load_examples = True|load_examples = False|g' airflow.cfg
grep dags_folder airflow.cfg
mkdir dags
```

```
cd dags  
ln -s $PWD/ml_pipeline_dag.py ~/airflow/dags
```



pyproject.toml:

```
version = "0.4"  
  
<...>  
  
"apache-airflow==2.10.5",
```

AULA 5.2 Criação de DAG e tasks



Debug console:

```
dvc_stages  
dvc_tasks  
airflow scheduler  
airflow webserver
```



dags/ml_pipeline_dag.py:

```
import yaml  
from pathlib import Path  
  
from airflow import DAG  
from airflow.operators.bash import BashOperator  
from airflow.operators.python import PythonOperator  
  
project_root = Path(__file__).resolve().parents[1]
```

```
def get_dvc_stages():
    dvc_yaml_path = project_root / "dvc.yaml"
    with open(dvc_yaml_path) as f:
        dvc_config = yaml.safe_load(f)
    return list(dvc_config["stages"].keys())

def register_artifacts_callable():
    from src.register_artifacts import main
    main()

default_args = {
    "owner": "airflow",
    "retries": 1,
}

with DAG(
    "ml_pipeline",
    default_args=default_args
) as dag:
    # DVC Pipeline Stages
    dvc_stages = get_dvc_stages()

    # Create tasks for each DVC stage
    dvc_tasks = []
    for stage in dvc_stages:
        task = BashOperator(
            task_id=f"dvc_{stage}",
            cwd=project_root,
            bash_command=f"dvc repro {stage}"
        )
        dvc_tasks.append(task)

    # Register artifacts in MLflow
    register_artifacts = PythonOperator(
        task_id="register_artifacts",
        python_callable=register_artifacts_callable
    )
```

```

# Deploy model by building and running Docker container
create_app_image = BashOperator(
    task_id="create_app_image",
    cwd=project_root,
    bash_command = "docker build -t ml-classifier ."
)

# Set dependencies
for i in range(len(dvc_tasks) - 1):
    dvc_tasks[i] >> dvc_tasks[i + 1]

# Connect the last DVC task to register_artifacts and then to create_app_image
dvc_tasks[-1] >> register_artifacts >> create_app_image

```

AULA 5.3 Dashboard

Comandos:

Terminal:

`docker image list`

AULA 5.4 Docker compose

Código:

`docker-compose.airflow.yaml`

```

services:
  postgres:
    image: postgres:13
    environment:
      - POSTGRES_USER=airflow
      - POSTGRES_PASSWORD=airflow
      - POSTGRES_DB=airflow
  networks:

```

```
- backend

dind:
  image: docker:24.0.5-dind
  privileged: true
  expose:
    - 2375
  environment:
    - DOCKER_TLS_CERTDIR=
  networks:
    - backend

airflow:
  build:
    context: .
    dockerfile: Dockerfile.airflow
  restart: always
  depends_on:
    - postgres
    - dind
  ports:
    - "8080:8080"
  environment:
    - AIRFLOW__CORE__EXECUTOR=LocalExecutor
    -
AIRFLOW__DATABASE__SQLALCHEMY_CONN=postgresql+psycopg2://airflow:airflow@postgres/airflow
    - AIRFLOW__CORE__LOAD_EXAMPLES=false
    - AIRFLOW__CORE__DAGS_FOLDER=/home/airflow/mlops_project/dags
    - DOCKER_HOST=tcp://dind:2375
    - DOCKER_HUB_USERNAME=${DOCKER_HUB_USERNAME}
    - DOCKER_HUB_TOKEN=${DOCKER_HUB_TOKEN}
    - DAGSHUB_USER_TOKEN=${DAGSHUB_USER_TOKEN}

  networks:
    - backend
  command: bash -c "
    airflow db init &&
    airflow users create --username admin --password admin --firstname admin --lastname
admin --role Admin --email admin@example.com &&
```

```
(airflow scheduler &) &&
airflow webserver
"
```

```
networks:
  backend:
```

Dockerfile.airflow:

```
FROM apache/airflow:2.10.5-python3.12

USER root

RUN groupadd airflow && usermod -aG airflow airflow

COPY . /home/airflow/mllops_project
RUN chown -R airflow:airflow /home/airflow/mllops_project

USER airflow

WORKDIR /home/airflow/mllops_project

RUN pip install .
RUN dvc init --no-scm
```

.env:

```
DOCKER_HUB_USERNAME=<your-username>
DOCKER_HUB_TOKEN=<your-token>
```

ml_pipeline_dag.py:

```
# bash_command = "docker build -t ml-classifier ."
bash_command=""""
docker build -t ${DOCKER_HUB_USERNAME}/ml-classifier .
echo ${DOCKER_HUB_TOKEN} | docker login -u ${DOCKER_HUB_USERNAME} --password-stdin
docker push ${DOCKER_HUB_USERNAME}/ml-classifier
"""
```

AULA 5.5 Airflow conteinerizado

Links:

Repositórios no DockerHub: <https://hub.docker.com/repositories/<your-username>>

Comandos:

Terminal:

```
docker compose -f docker-compose.airflow.yaml build  
docker image list  
docker compose -f docker-compose.airflow.yaml up  
docker run -p 5001:5001 <your-username>/ml-classifier
```